

FoodSense-MY · Dataset3 Interim v6 Results

YOLO11s capacity run (v6_s) · Validation-only · Not promoted (nano production path)

0.826

Best val mAP50-95

0.964

Best val mAP50

0.945

Best val Precision

0.918

Best val Recall

0.878

Val Mee Goreng recall


yolo11s

Model family


Headline checkpoints

Checkpoint	Precision	Recall	mAP50	mAP50-95
Best (epoch 78, selected)	0.945	0.918	0.964	0.826
Peak mAP50 (epoch 44)	0.925	0.923	0.965	0.813
Last (epoch 98)	0.924	0.931	0.964	0.820
v5 nano baseline (val)	0.894	0.899	0.945	0.793

Training curves (validation metrics by epoch)



Decision context and locked interim-v5 split



Key findings

Run B (yolo11s from COCO) completed on the locked interim-v5 split (batch 16, seed 42).

- Best val mAP50-95 0.826 at epoch 78
- Strongest Mee Goreng recall among interim runs (0.878)
- Beats v5 nano aggregate (0.793) on validation

Not used as production init: architecture mismatch with the YOLO11n deployment path. Nano follow-ups continued in interim v7 (freeze) → v8 (MG recovery).

Run A (v6_n_cos) was the schedule/fine-tune attempt; artifacts focus here on the completed v6_s capacity run.

